

6.780 Notes

Lecturer:

ANDREW LIU

Spring 2024

Last updated on Tuesday 13th February, 2024.

Contents

1	February 6, 2024	3
2	February 8, 2024	3
2.1	Bayesian Binary Hypothesis Testing	3
2.2	0-1 Loss	5
3	February 13, 2023	6
3.1	Review: Bayesian Hypothesis Testing	6
3.2	Non-Bayesian Hypothesis Testing	7
3.3	General Performance Measures	7
3.4	Operating Characteristic of the LRT (OC-LRT)	8
3.5	Neyman-Pearson Criterion	8

1 February 6, 2024

2 February 8, 2024

The setup for today's lecture is a model family $\mathcal{H} \in \{H_0, H_1, \dots, H_{M-1}\}$. In the classification problem, we can think of \mathcal{H} as a set of class labels, and we want to determine the correct label given some test data.

2.1 Bayesian Binary Hypothesis Testing

In this case, $M = 2$, so there are only two hypotheses. Our model has two major components. The first is some a priori information

$$P_0 = \mathbb{P}[H = H_0]$$

$$P_1 = \mathbb{P}[H = H_1] = 1 - P_0.$$

We also have the observation model, which is given by likelihood functions

$$H_0 : p_{Y|H}(\cdot|H_0)$$

$$H_1 : p_{Y|H}(\cdot|H_1).$$

Our goal is to create a **decision rule**, a.k.a. a **classifier**, which maps every $y \in \mathcal{Y}$ to some hypothesis $H_i \in \mathcal{H} = \{H_0, H_1\}$. This is somewhat confusing with the standard terminology of a hypothesis class being the set of possible solutions to a model, but we accept it for now.

Definition 2.1 (Cost)

In its most general form, we let

$$C(H_j, H_i) \triangleq C_{ij}$$

denote the cost of predicting H_j when the correct class is H_i .

Using cost to drive the notion of “best”, our best possible decision rule takes

the form

$$\hat{H}(\cdot) = \arg \min_f \mathbb{E}_{Y,H}[C(H, f(Y))].$$

The expected cost on the RHS is called **Bayes risk**, which we denote as $\varphi(f)$ for any decision rule f .

We can explicitly calculate this quantity:

$$\begin{aligned} \varphi(f) &= \mathbb{E}_{Y,H}[C(H, f(Y))] \\ &= \mathbb{E}_Y[\mathbb{E}_{H|Y}[C(H, f(Y))|Y = y]] \\ &= \int p_Y(y) \mathbb{E}[C(H, f(Y))|Y = y] dy. \end{aligned}$$

Notice that we have control over the expected risk for each point, so to minimize $\varphi(f)$, we only have to solve a solution for individual points. For a fixed $y^* \in \mathcal{Y}$, there are two possibilities; if $f(y^*) = H_0$, then

$$\mathbb{E}[C(H, f(y^*))|y = y^*] = C_{00}\mathbb{P}[H = H_0|y = y^*] + C_{01}\mathbb{P}[H = H_1|y = y^*],$$

otherwise

$$\mathbb{E}[C(H, f(y^*))|y = y^*] = C_{10}\mathbb{P}[H = H_0|y = y^*] + C_{11}\mathbb{P}[H = H_1|y = y^*].$$

This already technically gives us the optimal decision rule; for any given input y , we can explicitly compute both values, and return the hypothesis that gives the lesser of the two values. We can also express this in a simpler form. Since

$$\mathbb{P}[H = H_i|Y = y] = \frac{p_{Y|H}(y|H_i)p_H(H_i)}{p_Y(y)},$$

we can substitute into the above expressions:

$$C_{00}p_{Y|H}(y|H_0)P_0 + C_{01}p_{Y|H}(y|H_1)P_1 \stackrel{\hat{H}=H_1}{\geq} C_{10}p_{Y|H}(y|H_0)P_0 + C_{11}p_{Y|H}(y|H_1)P_1 \stackrel{\hat{H}=H_0}{\leq}$$

We can rewrite this expression in terms of the ratios

$$L(y) \triangleq \frac{p_{Y|H}(y|H_1)}{p_{Y|H}(y|H_0)} \stackrel{\hat{H}=H_1}{\geq} \frac{P_0(C_{10} - C_{00})}{P_1(C_{01} - C_{11})} \triangleq \eta.$$

We call $L(y)$ the **likelihood ratio**.

Theorem 2.2 (Likelihood Ratio Test)

Given a priori probabilities P_0, P_1 , data y , observation models $p_{Y|H}(\cdot|H_0), p_{Y|H}(\cdot|H_1)$, and costs $C_{00}, C_{01}, C_{10}, C_{11}$, the Bayesian decision rule form

$$L(y) \triangleq \frac{p_{Y|H}(y|H_1)}{p_{Y|H}(y|H_0)} \underset{\hat{H}_0}{\overset{\hat{H}_1}{\gtrless}} \frac{P_0(C_{10} - C_{00})}{P_1(C_{01} - C_{11})} \triangleq \eta,$$

meaning that the decision is $\hat{H}(y) = H_1$ when $L(y) > \eta$, $\hat{H}(y) = H_0$ when $L(y) < \eta$, and it is indifferent when $L(y) = \eta$.

Note that the optimal rule is simple and deterministic. Prof. Wornell makes a point about $L(y)$ being a scalar that we can always calculate. This is the heart of classification models; in larger neural nets, like ImageNet, ultimately what the large network of weights allows us to do is to express the intractable probabilities and compute a scalar value.

2.2 0-1 Loss

In the case of “0-1 loss”, i.e., $C_{00} = C_{11} = 0$, $C_{01} = C_{10} = 1$, in which case our test simplifies to

$$p_{H|Y}(H_1|y) \underset{\hat{H}_0}{\overset{\hat{H}_1}{\gtrless}} p_{H|Y}(H_0|y).$$

This is the **maximum a posteriori** (MAP) decision rule.

If we additionally assume that $P_0 = P_1$, i.e., that our prior belief is indifferent, then our test further simplifies to

$$p_{Y|H}(y|H_1) \underset{\hat{H}_0}{\overset{\hat{H}_1}{\gtrless}} p_{Y|H}(y|H_0).$$

This is the **maximum likelihood** (MLE) decision rule. In either case, the expected rate of error is given by

$$\varphi(\hat{H}) = \mathbb{P}[\hat{H}(Y) = H_0, H = H_1] + \mathbb{P}[\hat{H}(Y) = H_1, H = H_0].$$

Example 2.3 (Communicating a Bit)

We have a signal y , randomly distributed with variance σ^2 , and with two possible sources s_0, s_1 .

The likelihood ratio test gives

$$\ln L(y) = \ln \left(\frac{e^{-(y-s_1)^2/(2\sigma^2)}}{e^{-(y-s_0)^2/(2\sigma^2)}} \right) = \frac{1}{2\sigma^2} ((y-s_0)^2 - (y-s_1)^2).$$

Assuming 0-1 loss, $\ln L(y) = 0$, so the decision boundary is

$$y \underset{\hat{H}_0}{\overset{\hat{H}_1}{\gtrless}} \frac{s_0 + s_1}{2}.$$

We could compute the expected rate of error as follows:

$$\begin{aligned} \varphi(\hat{H}) &= \frac{1}{2} \left(\mathbb{P}[\hat{H}(Y) = H_0 | H = H_1] + \mathbb{P}[\hat{H}(Y) = H_1 | H = H_0] \right) \\ &= \frac{1}{2} \left(\mathbb{P} \left[y < \frac{s_0 + s_1}{2} \middle| H = H_1 \right] + \mathbb{P} \left[y \geq \frac{s_0 + s_1}{2} \middle| H = H_0 \right] \right) \\ &= \frac{1}{2} \left(\mathbb{P} \left[\frac{y - s_1}{\sigma} < \frac{s_0 - s_1}{2\sigma} \middle| H = H_1 \right] + \mathbb{P} \left[\frac{y - s_0}{\sigma} \geq \frac{s_1 - s_0}{2\sigma} \middle| H = H_0 \right] \right) \\ &= Q \left(\frac{s_1 - s_0}{2\sigma} \right). \end{aligned}$$

The quantity $(s_1 - s_0)/\sigma$ is a measure of signal-to-noise; the larger the SNR, the more uncertain we are about our prediction, i.e., the higher our expected rate of error.

3 February 13, 2023

Bad weather day today. Prof. Wornell tells us about how this is the first time he's given a lecture over zoom since the pandemic. He was hoping that he would never have had to give a zoom lecture again, but here we are.

3.1 Review: Bayesian Hypothesis Testing

Overarching goal: optimize the expected cost of choosing one hypothesis over another. This can be accomplished with the Likelihood Ratio Test:

$$L(y) \underset{\hat{H}_0}{\overset{\hat{H}_1}{\gtrless}} \eta$$

Some issues with this solution: we require some notion of costs, as well as priors $P_i = P_H(H_i)$. It can be difficult to assign probabilities to abstract concepts, like $\mathbb{P}[\text{patient has disease}]$.

3.2 Non-Bayesian Hypothesis Testing

The “folk theorem” that we will be proving today is that all optimum decision rules takes the form of a Likelihood Ratio Test:

$$\frac{p_{Y|H}(y|H_1)}{p_{Y|H}(y|H_0)} \underset{\hat{H}_0}{\overset{\hat{H}_1}{\gtrless}} \eta,$$

for some η . This is largely true, but not always.

3.3 General Performance Measures

Let $\hat{H}(\cdot)$ be any rule. An equivalent characterization of this rule is some partition of the observation space Y :

$$\begin{aligned} y_0 &= \{y \in Y : \hat{H}(y) = H_0\} \\ y_1 &= Y \setminus y_0 \end{aligned}$$

Then $P_D = \mathbb{P}[\hat{H} = H_1 | H = H_1] = \int_{y_1} p_{Y|H}(y|H_1) dy$ is called the **detection probability**, and $P_F = \mathbb{P}[\hat{H} = H_1 | H = H_0] = \int_{y_1} p_{Y|H}(y|H_0) dy$ is called the **“false alarm” probability**. Some related terminology: $P_M = 1 - P_D$ is called the “miss” probability.

- Statistics terminology: $P_E^1 = P_F$, $P_E^2 = P_M$, “probability of error of each kind”. The probability of error of the first kind is called the “size” of the test, while the probability of error the second kind is called the “power” of the test.
- Medical terminology: P_F is the false positive rate, while P_M is the false negative rate.
- Learning / pattern classification: $P_R = P_D$ is the “recall” or “sensitivity”. The “precision” is defined as $P_P = \mathbb{P}[H = H_1 | \hat{H} = H_1] = 1/(1 + P_F/P_D \cdot P_0/P_1)$.

In general, P_D and P_F are conflicting objectives. We seek large P_D and small P_F . The bayesian approach to this “multi-objective” optimization is to choose the rule that satisfies:

$$\min_{\hat{H}(\cdot)} (\alpha P_F - \beta P_D).$$

3.4 Operating Characteristic of the LRT (OC-LRT)

Other tradeoffs are possible. Consider the family of ratio tests:

$$\{\hat{H}(\cdot) = \text{LRT, for some } \eta\}.$$

Example 3.1

Consider two hypotheses

$$H_0 : y \sim \mathcal{N}(0, \sigma^2)$$

$$H_1 : y \sim \mathcal{N}(m, \sigma^2)$$

The LRT is given by

$$y \underset{\hat{H}_0}{\overset{\hat{H}_1}{\gtrless}} \frac{m}{2} + \frac{\sigma^2 \ln \eta}{m} \triangleq \gamma.$$

add images!

This curve is called the **OC-LRT**: Operating Characteristic of the Likelihood Ratio Test.

$$\text{OC-LRT} : \{(P_F, P_D) : \hat{H}(\cdot) \text{ is LRT for some } \eta\}.$$

Claim 3.2

OC-LRT is monotonic and non-decreasing.

Proof. If $\eta_2 > \eta_1$, then $P_D(\eta_2) \leq P_D(\eta_1)$ and $P_F(\eta_2) \leq P_F(\eta_1)$. □

3.5 Neyman-Pearson Criterion

To avoid the problem of costs and priors, choose a rule subject to

$$\max_{\hat{H}(\cdot)} P_D \text{ s.t. } P_F \leq \alpha.$$

Theorem 3.3 (Neyman-Pearson Lemma, Special Case)

For deterministic $\hat{H}(\cdot)$, the solution to the NP-Criterion is an LRT when the LRT is continuous. In other words,

$$\hat{H}(y) = H_{\mathbb{1}_{L(y) \geq \eta}},$$

where η is the smallest threshold s.t.

$$P_F = \mathbb{P}(L(y) \geq \eta | H = H_0) \leq \alpha.$$

Proof. We can prove this with lagrange multipliers. Fix $P_F = \alpha' \leq \alpha$. Then, we want to optimize

$$\begin{aligned} \min_{\hat{H}(\cdot)} \varphi(\hat{H}) &= (1 - P_D) + \lambda(P_F - \alpha') \\ &= \int_{y_0} p_{Y|H}(y|H_1) dy + \lambda \left(\int_{y_1} p_{Y|H}(y|H_0) dy - \alpha' \right) \\ &= \lambda(1 - \alpha') + \int_{y_0} (p_{Y|H}(y|H_1) - \lambda p_{Y|H}(y|H_0)) dy. \end{aligned}$$

The min φ occurs when we assign y to y_0 whenever the the integrand is ≤ 0 , as this minimizes the cost. For the same reason, we want α' to be as large as possible. Therefore,

$$\frac{p_{Y|H}(y|H_1)}{p_{Y|H}(y|H_0)} \underset{\hat{H}_0}{\overset{\hat{H}_1}{\gtrless}} \lambda,$$

where α' is chosen to be the largest threshold achievable by LRT. □

This falls apart slightly when $L(y)$ has discrete components.

Example 3.4 (OC-LRT falls apart under discrete likelihood)

draw some graphs.